

CS224N: NATURAL LANGUAGE PROCESSING WITH DEEP LEARNING

ASSIGNMENT #3

ANTHONY HO

1. (a) (i) Example 1: “Stanford is great.” - where “Stanford” could refer to Stanford University (organization) or a person with last name Stanford (person).
Example 2: “I am going to Stanford.” - where “Stanford” could refer to Stanford University (organization) or Stanford, California (location).
- (ii) Because the features apart from the word itself could provide additional information and context not contained in the word itself which might help with removing ambiguity in identifying its named entity.
- (iii) Feature 1: the adjacent words from the word itself could be helpful in predicting whether the word is part of a named entity or not. For example, if the word is immediately succeeded by an action verb, it makes the word more likely to be a named entity.
Feature 2: capitalization of the word could also be helpful in predicting whether the word is part of a named entity or not, especially in the case of person, organization, and location.

- (b) (i) The dimensions are:

$$\mathbf{e}^{(t)} \in \mathbb{R}^{1 \times D(2w+1)}$$

$$\mathbf{W} \in \mathbb{R}^{D(2w+1) \times H}$$

$$\mathbf{U} \in \mathbb{R}^{H \times C}$$

- (ii) The computational complexity of predicting labels for a single word is $\mathcal{O}(D(2w+1) + D(2w+1)H + HC) = \mathcal{O}(D(2w+1)(H+1) + HC)$. Therefore the computational complexity of predicting labels for a sentence of length T is $\mathcal{O}(T(D(2w+1)(H+1) + HC))$.
- (c) Please see the coding portion of the assignment.
- (d) (i) The best development entity-level F_1 score is 0.83 and the corresponding token-level confusion matrix is shown below:

| | PER | ORG | LOC | MISC | O |
|------|---------|---------|---------|--------|----------|
| PER | 2958.00 | 51.00 | 61.00 | 10.00 | 69.00 |
| ORG | 139.00 | 1659.00 | 120.00 | 47.00 | 127.00 |
| LOC | 48.00 | 148.00 | 1843.00 | 19.00 | 36.00 |
| MISC | 44.00 | 71.00 | 46.00 | 995.00 | 112.00 |
| O | 49.00 | 50.00 | 17.00 | 30.00 | 42613.00 |

From the confusion matrix, it looks like the model has a tendency to misclassify organization as person, location or null, to misclassify location as organization, and to misclassify miscellaneous as null.

- (ii) (1) A window-based model would have troubles identifying named entities longer than the window itself. For example, the prediction made by our model (`window_size = 1`) on the sentence “The Senate Select Committee on Intelligence is investigating the Russian affairs .” is “O ORG ORG ORG O ORG O O O MISC O O”, which fails to identify the “Senate Select Committee on Intelligence” as a single named entity instead of two.
 - (2) A window-based model would fail to take long-range information into account, since it’s based on a finite length window. For example, the prediction made by our model (`window_size = 1`) on the sentence “Washington was the first President of the United States .” is “LOC O O O O O O LOC LOC O”, which fails to take into the account of the word “President” that establishes “Washington” as a person instead of a location.
2. (a) (i) This particular RNN model does not necessarily have more parameters in comparison to the window-based model, since the window-based model has a bigger W (due to a bigger window) than the RNN’s

model's W_e , while the RNN model has the additional parameter W_h . The difference in number of parameters between the RNN model over the window-based model is $H^2 - 2wDH$.

- (ii) The computational complexity mainly depends on the sizes of the matrix multiplications and scales linearly to the number of time steps. Thus, the computational complexity of predicting labels for a sentence of length T with the RNN model is $\mathcal{O}(T(H^2 + DH + CH))$.
- (b) (i) Since the cross-entropy cost is computed at the token level, its minimization does not guarantee the minimization of the entity-level F_1 score, such as when there are multi-word entities in the training data.
- (ii) F_1 score is not additive over training samples like cross-entropy cost, which makes it difficult to be used in stochastic gradient descent during optimization (e.g. can't compute the F_1 score of the full training data from the F_1 scores of the individual minibatches).
- (c) Please see the coding portion of the assignment.
- (d) (i) If we did not use masking, the loss would be contaminated by the cross-entropy cost associated with the NULL tokens and the zero label, which effectively decreases the signal-to-noise of the loss function. It could negatively affect learning, especially when the maximum sentence length is far greater than the average sentence length.
- (ii) Please see the coding portion of the assignment.
- (e) Please see the coding portion of the assignment.
- (f) Please see the coding portion of the assignment.

- (g) (i) (1) This RNN model could suffer from the vanishing gradient problem, which means that long-range information might not be taken into account. An example of this from the model's output is shown below (truncated since it's too long to fit):

```
x : Automakers petitioned the National Highway Traffic Safety Administration to allow ...
y*: 0          0          0  ORG      ORG      ORG      ORG      ORG          0  0      ...
y': 0          0          0  ORG      ORG      ORG      ORG      0          0  0      ...
```

We can see that the model fails to consider "Administration" as part of a very long organization name, which suggests the vanishing gradient problem is an issue with this model.

(2) This RNN model is unidirectional, which means information that happens after a word is not taken into account. An example of this from the model's output is shown below (truncated since it's too long to fit):

```
x : Jordan , whose 1992 film " The Crying Game " also came under fire for what was ...
y*: PER    0 0      0    0    0  MISC MISC  MISC 0 0      0    0      0    0      0    0      ...
y': LOC    0 0      0    0    0  0    MISC MISC 0 0      0    0      0    0      0    0      ...
```

We can see that the model fails to classify "Jordan" as a person, even though it is immediately succeeded by "whose", which highly suggests "Jordan" is a person. This seems to be especially a problem with the first few words in a sentence.

- (ii) (1) LSTM or GRU
 - (2) Bidirectional RNN
3. (a) (i) A combination of values of w_h , u_h and b_h for an RNN cell is:

$$\begin{aligned} w_h &= 1 \\ u_h &= 1 \\ b_h &= 0 \end{aligned}$$

- (ii)
- (b) (i)
- (ii)

- (c) Please see the coding portion of the assignment.
- (d) Please see the coding portion of the assignment for implementation. The plots of learning dynamics are shown in figure (1)–(4).




FIGURE 1. Plots of learning dynamics generated for a RNN without gradient clipping.




FIGURE 2. Plots of learning dynamics generated for a RNN with gradient clipping.

- (e) (i) Both RNN models experience vanishing gradients and plateau at a final loss of 12.4978. Gradient clipping does not help with the vanishing gradients problem. Both GRU models experience exploding gradients at around timestep 1100, but in this case gradient clipping helps bring the problem under control.
- (ii) The GRU model with gradient clipping performs the best, arriving at a final loss of 8.0668. It is likely because the model is equipped to handle both vanishing gradients (GRU) and exploding gradients (gradient clipping) problems.

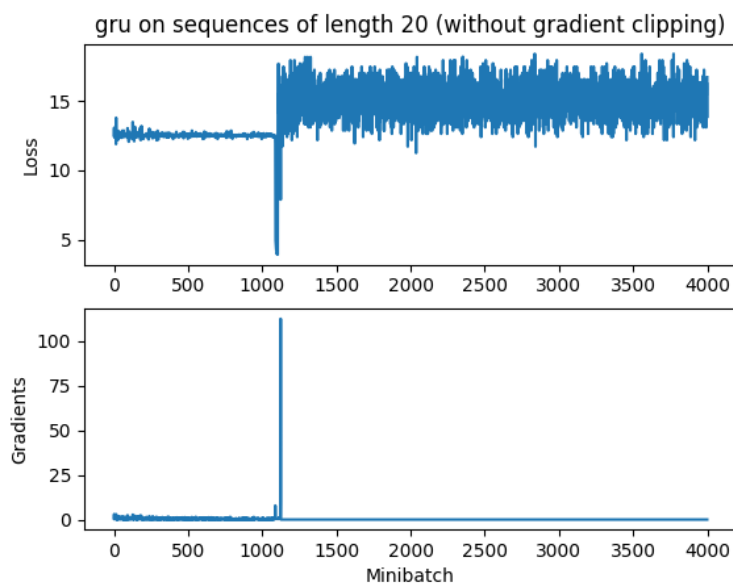


FIGURE 3. Plots of learning dynamics generated for a GRU without gradient clipping.




FIGURE 4. Plots of learning dynamics generated for a GRU with gradient clipping.

(f) Please see the coding portion of the assignment.